

Now cuts hook overhead 85%

ONE AGENT TO ROUTE THEM ALL

The installer runs end-to-end. Prompt 0 picks the orchestrator (claude, cursor, cline, codex, aider or generic). Prompt 7 collects API keys with per-provider instructions and writes a `0600 .env`. Setup now also writes a `CLAUDE.md` doctrine file that is auto-loaded and prompt-cache friendly, so the routing map, keep-in-Claude rules, and 4-step pipeline never re-inject on every message. Multi-framework users get a lean sibling-detected block that skips shared sections. Measured: **85% cut in hook overhead** vs the prior version on a 50-message session.


What changed since the last PDF: SessionStart shrank from ~1500 to ~250 bytes (routing map moved to `CLAUDE.md`). UserPromptSubmit dropped from ~820 to ~70 bytes per message. `CLAUDE.md` carries the doctrine once and Claude Code caches it for one hour. The `pal-router` skill gained 11 token-efficiency rules (auto-compress outputs over 5 KB, batch PAL calls, { claim, confidence, source_span } triples, in-conversation failure-map, diff-first, response terseness ladder). Sibling detection dedups doctrine when both sauron and let-there-be-light are installed.

1. The orchestrator choice

c	claude-code (default)	writes <code>~/.claude/settings.json</code> hooks + symlinks skills into <code>~/.claude/skills/</code>
r	cursor	writes <code>./.cursor/rules/sauron.mdc</code> with <code>alwaysApply:true</code> ; copies skills alongside
l	cline	writes <code>./.clinerules</code> (or <code>~/.clinerules</code>); copies skills alongside
x	codex-cli	writes <code>~/.codex/instructions.md</code> ; copies skills into <code>~/.codex/sauron-skills/</code>
a	aider	writes <code>./.aider.sauron.md</code> ; add to <code>.aider.conf.yml</code> as read:
g	generic	writes <code>./SYSTEM_PROMPT.sauron.md</code> you paste into any tool; copies skills alongside

2. The terminal session (Claude Code path)

~.~.~.~.
.(0).
~.~.~.~.



interactive setup . one agent to route them all

0. Which agent orchestrator do you use?

- c) claude-code (default) full: hooks + Skill() invocations + delegate policy
 - r) cursor writes .cursor/rules/Framework.mdc (alwaysApply)
 - l) cline writes .clinerules with the delegate policy
 - x) codex-cli (OpenAI) writes ~/.codex/instructions.md
 - a) aider writes .aider.conf.yml + Framework.md conventions
 - g) generic / other writes SYSTEM_PROMPT.md you paste into any tool
- ✓ orchestrator: c

1. Where should it install?

- g) global → ~/.claude/settings.json (loads every session, every project)
- p) per-project → \$PWD/.claude/settings.json (only loads when Claude Code runs here)
- s) skip → print the settings JSON to stdout, do not write anything

2. Which skills should auto-load at session start?

These are invoked as the very first tool calls of every session.
Leave blank to skip a skill; type y to include it.

3. Which PAL models should be listed in the delegate-first policy?

Pick every model your PAL install actually has keys for.
Unselected models are dropped from the routing map so Claude never tries them.

Ready to write /tmp/sauron-pdf-v4/testhome/.claude/settings.json

Existing file will be backed up with .bak.<timestamp> suffix.

- ✓ wrote /tmp/sauron-pdf-v4/testhome/.claude/settings.json
- ✓ wrote CLAUDE.md doctrine: /tmp/sauron-pdf-v4/testhome/CLAUDE.md
- ! skipped skill symlinking; you must place skills under /tmp/sauron-pdf-v4/testhome/.claude/skills manually
- ! ~/.claude.json not found; make sure PAL is registered before starting Claude Code
- ✓ wrote env template: /tmp/sauron-pdf-v4/testhome/.claude/.env.sauron.example

7. Enter API keys now to complete installation?

You can skip and edit .env.sauron.example later, or enter them now to have a real .env.sauron written with 0600 permissions.
skipped interactive key entry; only the example was written.
! edit /tmp/sauron-pdf-v4/testhome/.claude/.env.sauron.example and rename to .env.sauron before starting your orchestrator.

Next steps

1. Register PAL as an MCP server in ~/.claude.json:
"mcpServers": { "pal": { "type": "stdio", "command": "/path/to/zen-mcp-server/.pal_venv/bin/python", "args": ["/path/to/zen-mcp-server/server.py"], "env": { ...keys... } } }
2. Source your API keys: `source /tmp/sauron-pdf-v4/testhome/.claude/.env.sauron` (after editing it)
3. Restart Claude Code.
4. On the next session start you should see the auto-invoked skills fire immediately.

✓ sauron setup complete.

3. What Claude sees at the next session start (now lean)

The `SessionStart` hook fires *before* Claude responds. Its payload is injected into context. The full doctrine now lives in `CLAUDE.md`, so the hook is just a pointer.

SessionStart · additionalContext (163 bytes)

Doctrine + routing map are in `CLAUDE.md` at the project root (already loaded). Preload skills now: ``caveman`` (full mode), ``pentesting-agent``, ``validator`` (preload).

UserPromptSubmit · additionalContext (69 bytes, every message)

Doctrine active from `CLAUDE.md`. Delegate + failover. Batch PAL calls.

Why it is cheap: the routing map, keep-in-Claude list, failover doctrine, 4-step pipeline, and the 11 token-efficiency rules moved into `CLAUDE.md`. Claude Code auto-loads that file into the system prompt and caches it for one hour. On message 2 through N in the same session you pay 0 tokens for it. Prior version re-injected all of it on every user message.

4. The `CLAUDE.md` doctrine file (4390 bytes, loaded once)

Setup writes this alongside `settings.json`. It sits inside a `<!-- BEGIN sauron doctrine -->` managed block so re-running `setup.sh` is idempotent. When both `sauron` and `let-there-be-light` are installed, the second one detects its sibling and writes a lean block that skips the shared sections.

```

<!-- BEGIN sauron doctrine (managed by setup.sh; regenerate to update) -->
# sauron doctrine

## Delegate-first routing
groq (gpt-oss-120b, ~500 rpm, 8000 tpm cap) for report writing, vuln explanations, and
skeptical validation. nemotron (nvidia, 1M ctx) for bulk reading of large files. grok (x-ai,
2M ctx) for permissive high-context security reasoning when Gemini refuses. flash
(gemini-3.6-flash, 1M ctx) for fast structured extraction. or-free (OpenRouter meta-router,
200K ctx) as generalist fallback. pro (gemini-3-pro-preview) for deep-reasoning fallback and
adversarial debate via mcp_pal_challenge.

## Keep in Claude ONLY
- user-facing decisions
- exploitation choices, severity calls
- safety-boundary checks (no 3rd-party data, no destructive actions, no account creation)
- side-effecting actions
- tool-sequence orchestration

## Failover
Any model refusal or error routes to the next model in the map. A classifier refusal is a
routing problem, not a stop.

## 4-step validation pipeline (per confirmed finding, before presenting or reporting)
A. Stress-test via validator skill (false-positive check, reproducibility, evidence gaps,
confidence).
B. Run the gap tests, negative controls, and live observations the validator flags.
C. Adversarial challenge via mcp_pal_challenge (pro primary, groq fallback): challenge
severity, exploitability, chain viability, impact ceiling.
D. Synthesize with confidence markers; groq drafts the report, Claude spot-checks low-
confidence claims only.

## Token efficiency rules (this framework exists to cut Claude token cost)
- Any tool output over 5 KB routes through nemotron (bulk) or flash (structured) for
summarization before Claude reads.
- Batch related PAL sub-tasks into one structured call, not N separate ones.
- PAL responses on factual output must include {claim, confidence, source_span} triples.
Claude byte-checks entries below high confidence only.
- Never re-Read a file already Read this session; recall from conversation context.
- Preemptively bound predictable-bloat outputs: append | head -c 5000 or | jq -c on Bash
calls whose full output you don't need.
- Never delegate the same task twice; cache the result inline in the conversation.
- Terser confirmations when a diff or file speaks for itself; not every action needs a
paragraph.

## Response terseness ladder
- Level 1 (one-liner): use when a diff, SHA, file i

... [truncated for PDF; full file is 4390 bytes]

```

5. Files on disk after install

```

.claude
.claude/.env.sauron.example

```

CLAUDE.md
.claude/settings.json

.env.sauron.example (only the API keys your model picks need)

```
# sauron API keys - source this from your shell rc, or export before starting Claude Code.  
# Do NOT commit the real values.  
export GEMINI_API_KEY=your-gemini-key  
export OPENROUTER_API_KEY=your-openrouter-key  
export CUSTOM_API_URL=https://api.groq.com/openai/v1  
export CUSTOM_API_KEY=your-groq-key
```

Repo: github.com/crowx01/sauron · Sibling: [crowx01/let-there-be-light](https://github.com/crowx01/let-there-be-light) (dev-lifecycle)